

Evaluation Questions

1. You are required to create a Python dictionary with two keys: data and target. The data key should store a 28x28 matrix filled with random integer values between 0 and 255 (both values included), representing pixel intensities in a grayscale image. The target key should store a binary label (either 0 or 1, chosen randomly). Then print the shape and type of the dictionary elements.

(25 mins)

Steps:

- i. Use the randint function from Python's random module to generate the 28x28 matrix of random integer values ranging between 0 and 255, both ends included. You can use internet to look at documentation of [randint](#), and not for any other purpose.
- ii. Convert the data generated in step 1 to a numpy array.
- iii. Assign a binary label (either 0 or 1, chosen randomly) as the target.
- iv. Construct the dictionary with the following structure:

data: The 28x28 matrix generated in step 1.

target: label obtained in step 2.

- v. Print the shape and type of dictionary elements.

2. You can go through the [logistic regression docs](#). Then do the following.

(1 hr)

[Note: For this exercise, you can continue from the end in the same notebook you submitted; you can reuse the code you have already written wherever required.]

- i. Implement a logistic regression model for 5 versus non-5 mnist digit classification wherein the data imbalance must be handled automatically. Keep a hold out test set of 200 samples wherein 100 samples are from class 5 and 100 samples are from non-5 classes. The choice of choosing the 100 samples from each class is left to you.
- ii. Record the average train and test set accuracy.
- iii. Incorporate L1 regularization to the above model. Choose the regularization parameter based on the results in step (ii).
- iv. Record the average train and test set accuracy.
- v. Incorporate L2 regularization to the above model. The regularization parameter chosen in step (iii) can be used here.
- vi. Record the average train and test set accuracy.
- vii. Print the number of parameters of the model strictly less than $1e-5$ in both L1 and L2 regularization implementations.

3. Speak for 2 minutes on your attempt in question 2.

(2 mins)